
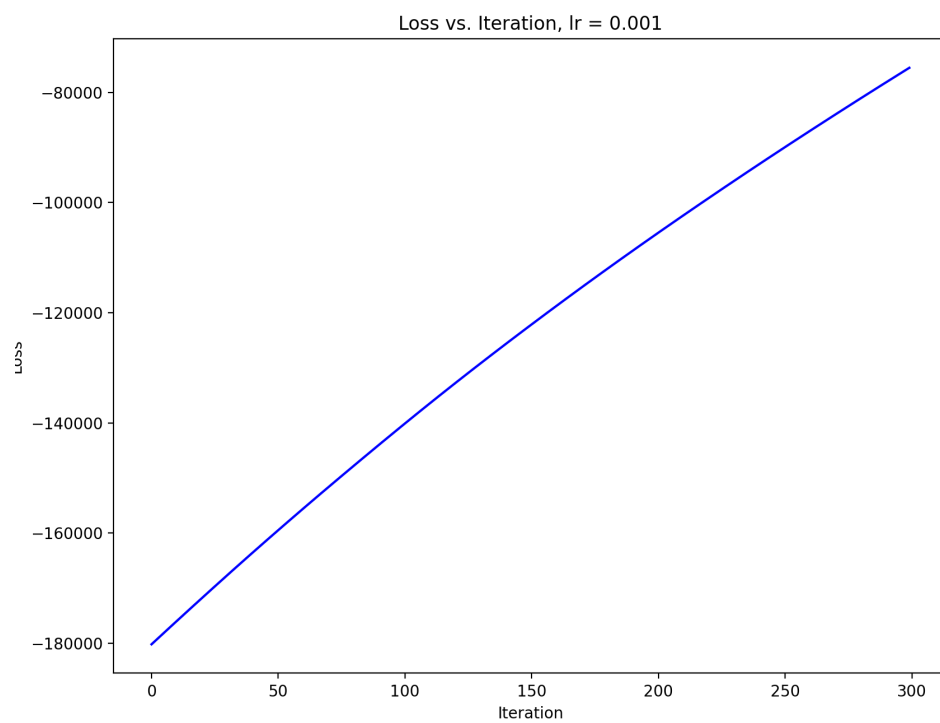





**Jeret McCoy, Abelardo Riojas, Gustavo Flores**  
**Part A**

LearnRate	Train Error	Test Error
0.01	11.7%	25.7%
0.05	50%	50%
0.001	40.6%	48.5%
0.005	20.6%	23.3%
0.0075	14.9%	19.2%









```
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd

class LiteLR:
    def __init__(self, learnRate, X, y, iterations=300, shrinkage=0.0001):
        self.iterations = iterations
        self.shrinkage = shrinkage
        self.learnRate = learnRate
        self.X = X
        self.y = y
        self.weights = np.random.randn(X.shape[1])
    def L(self):
        loss = 0
        for j, obj in enumerate(self.X): #for each row in X
```

```

        loss = loss + ( (self.y[j] * np.dot(self.weights, obj)) -
np.log(1 + np.exp( np.dot(self.weights, obj)  )) )
    return loss
def dL(self):
    dloss = 0
    for j, obj in enumerate(self.X): #for each row in X
        dloss = dloss + ( obj * (self.y[j] - (
(np.exp(np.dot(self.weights,obj)))/(1 + np.exp(np.dot(self.weights,obj)))
) ) )
    return dloss
def gradientAscent(self, plot=False): #trains model, returns loss
    losses = list()
    for i in range(self.iterations):
        self.weights = self.weights - (self.learnRate * self.shrinkage
* self.weights) + ((self.learnRate/self.X.shape[1])* self.dL())
        losses.append(self.L())
    if (plot == True):
        itrs = range(self.iterations)
        plt.plot(itrs, losses , c='b')
        plt.xlabel("Iteration")
        plt.ylabel("Loss")
        plt.title("Loss vs. Iteration, lr =
{}".format(self.learnRate))
        plt.show()
    return self.L()
def error(self, X, y):
    ret = 0
    for j, obj in enumerate(X):
        hyp = 1 / (1 + np.exp(-1 * np.dot(self.weights, obj)))
        pred = -1
        if (hyp > .5):
            pred = 1
        if (pred != y[j]):
            ret = ret + 1
    return ret / X.shape[0]

trainX = pd.read_csv("gisette_train.data", sep=' ', header=None)
trainX = np.array(trainX)
trainX = trainX[:, 0:5000]

```

```

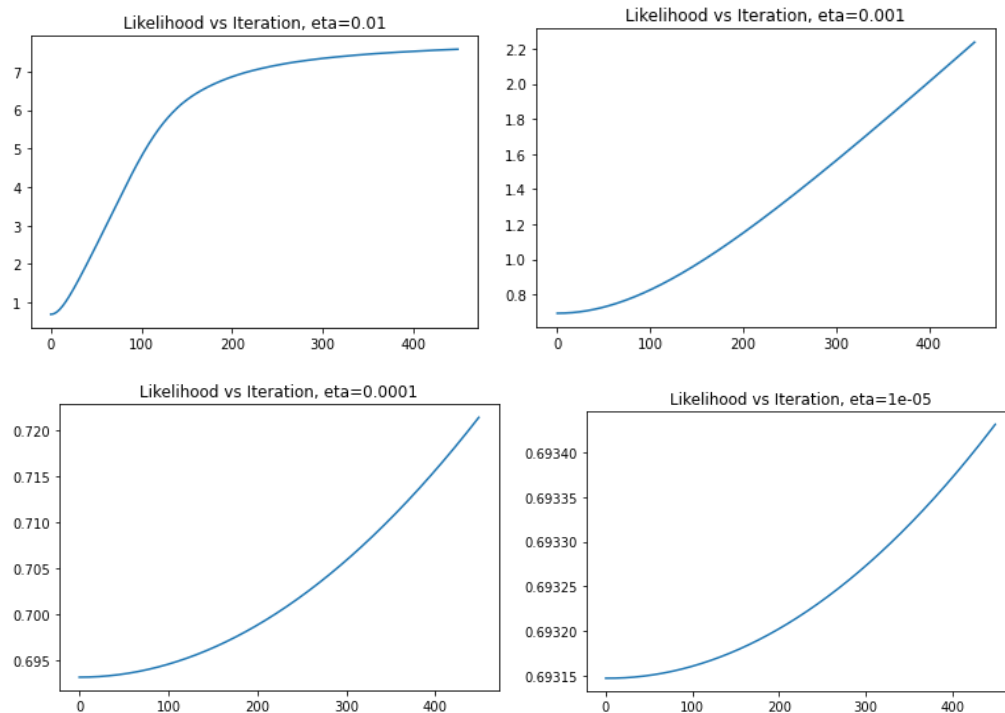
testX = pd.read_csv("gisette_valid.data", sep=' ', header=None)
testX = np.array(testX)
testX = testX[:, 0:5000]
trainy = pd.read_csv("gisette_train.labels", sep=' ', header=None)
trainy = np.array(trainy)
testy = pd.read_csv("gisette_valid.labels", sep=' ', header=None)
testy = np.array(testy)
for i in range(5000):
    if (np.std(trainX[:,i]) != 0):
        trainX[:,i] = (trainX[:,i] - np.mean(trainX[:,i])) /
np.std(trainX[:,i])
        testX[:,i] = (testX[:,i] - np.mean(trainX[:,i])) /
np.std(trainX[:,i])

clf = LiteLR(.0075, trainX, trainy)
loss = clf.gradientAscent(plot=True)
print(loss)
print("Train error: {} Test error: {}".format(    clf.error(trainX,
trainy) , clf.error(testX,testy)    ))

```

Part B:

Learning Rate	Train Misclass Error	Test Misclass Error
.01	15%	15%
.001	15%	15%
.0001	15%	15%
.00001	15%	15%



```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

X = pd.read_csv('X.dat', sep = ' ', header=None)
Xtest = pd.read_csv('Xtest.dat', sep = ' ', header=None)
Y = pd.read_csv('Y.dat', sep = ' ', header=None)
Ytest = pd.read_csv('Ytest.dat', sep = ' ', header=None)

for i, obs in enumerate(X):
    X[i] = (obs - obs.mean())/(obs.std())
for i, obs in enumerate(Xtest):
    Xtest[i] = (obs - obs.mean())/(obs.std())

class MyLogisticRegression:
    def __init__(self, learning_rate=0.01, max_iterations=450,
shrinkage=.0001):

        # Initialising all the parameters
        self.learning_rate = learning_rate
        self.max_iterations = max_iterations
        self.shrinkage = shrinkage
        self.likelihoods = []
```

```

        # Define epsilon because log(0) is not defined
        self.eps = 1e-7

    def sigmoid(self, z):

        ### START CODE HERE
        sig_z = (1/(1+np.exp(-z)))
        ### END CODE HERE

        assert (z.shape==sig_z.shape), 'Error in sigmoid implementation. Check
carefully'
        return sig_z

    def log_likelihood(self, y_true, y_pred):

        # Fix 0/1 values in y_pred so that log is not undefined
        y_pred = np.maximum(np.full(y_pred.shape, self.eps),
np.minimum(np.full(y_pred.shape, 1-self.eps), y_pred))

        likelihood = (y_true*np.log(y_pred)+(1-y_true)*np.log(1-y_pred))
        #self.likelihoods.append(-1*np.mean(likelihood))

        return np.mean(likelihood)

    def fit(self, X, y):

        num_examples = X.shape[0]
        num_features = X.shape[1]

        ### START CODE HERE

        # Initialize weights with appropriate shape
        self.weights = np.zeros((X.shape[1]))
        #print("Z",self.weights.shape)
        # print(X.shape)

        # Perform gradient ascent
        for i in range(self.max_iterations):

```



```

        z = np.dot(X,self.weights)

        # Output probability value by appplying sigmoid on z
        y_pred = self.sigmoid(z)

        # Calculate the gradient values
        # This is just vectorized efficient way of implementing gradient.
        gradient = np.mean(np.dot(X.T, y-y_pred), axis=1)
        #print(gradient.shape)

        # Update the weights
        # It is gradient ASCENT not descent
        self.weights = (self.weights) -
        (self.learning_rate*self.shrinkage*self.weights) +
        (self.learning_rate/X.shape[1])*gradient

        # Calculating log likelihood
        likelihood = self.log_likelihood(y,y_pred)

        self.likelihoods.append(likelihood)

#### END CODE HERE

def predict_proba(self,X):

    if self.weights is None:
        raise Exception("Fit the model before prediction")

    #### START CODE HERE

    z = np.dot(X,self.weights)
    probabilities = self.sigmoid(z)
    # probabilities.reshape(probabilities.shape[0],1)

    #### END CODE HERE

    return probabilities

def predict(self, X, threshold=0.5):

```

```

        # Thresholding probability to predict binary values

        binary_predictions = np.array(list(map(lambda x: 1 if x>threshold else
0, self.predict_proba(X))))

        return binary_predictions

clf = MyLogisticRegression(learning_rate=.01)

clf.fit(X,Y)

preds = clf.predict(X)

count_train = 0
for pred, y in zip(preds, Y):
    if pred == y:
        #print(pred, y)
        count_train+=1

preds = clf.predict(Xtest)

count_test = 0
for pred, y in zip(preds, Ytest):
    if pred == y:
        count_test+=1

print('Train Error: {:.2f}, Test Error
 {:.2f}'.format(1-count_train/X.shape[0], 1-count_test/Xtest.shape[0]))

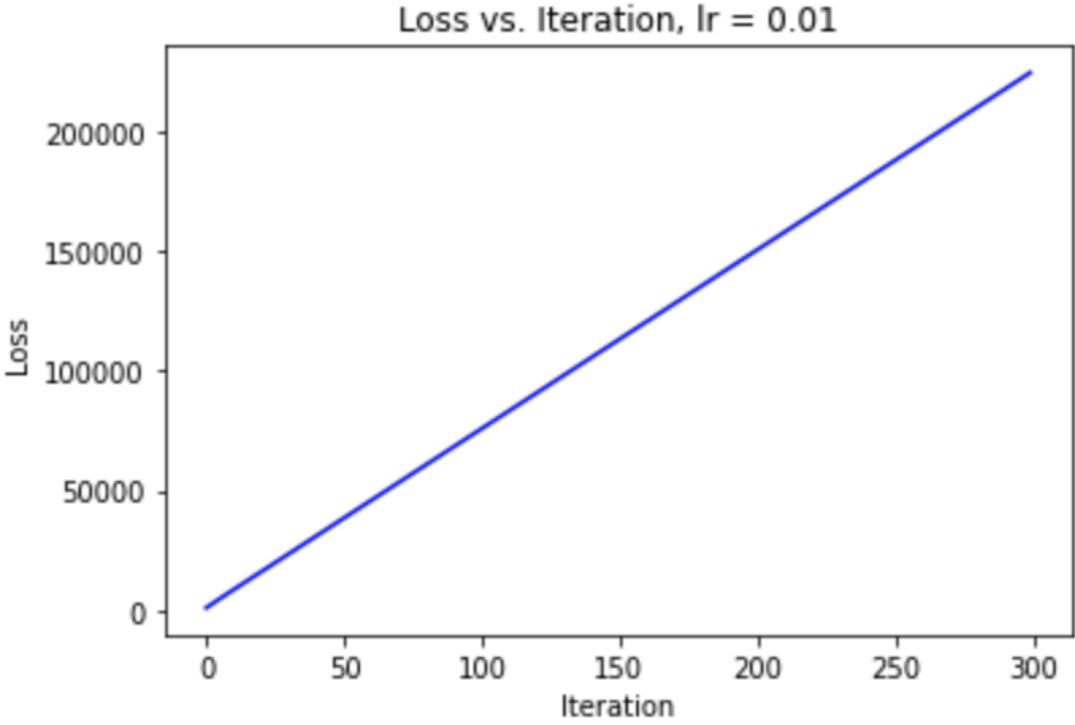
plt.title('Likelihood vs Iteration, eta={}'.format(clf.learning_rate))
plt.plot(np.abs(clf.likelihoods))
plt.show()

```


Part C.

LearnRate	Train Error	Test Error
0.01	44.8%	47.5%
0.05	49.2%	49.8%


0.001	44.8%	47.8%
0.0005	44.8%	48.5%
0.00001	49.8%	49.8%

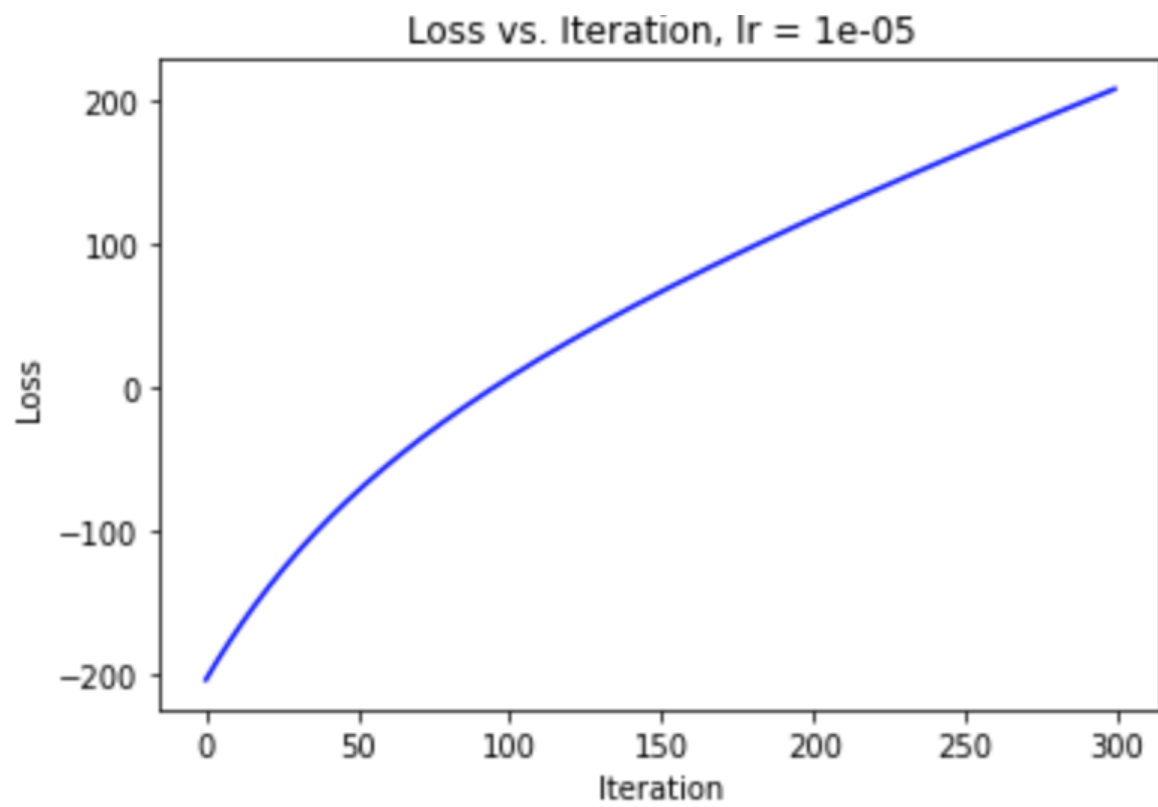



Loss vs. Iteration, lr = 0.05



Loss vs. Iteration, lr = 0.001





```
import numpy as np
import matplotlib.pyplot as plt
```

```

import pandas as pd
import csv

class LiteLR:
    def __init__(self, learnRate, X, y, iterations=300, shrinkage=0.0001):
        self.iterations = iterations
        self.shrinkage = shrinkage
        self.learnRate = learnRate
        self.X = X
        self.y = y
        self.weights = np.zeros(X.shape[1])
        #display(self.weights)
    def L(self):
        loss = 0
        for j, obj in enumerate(self.X): #for each row in X
            loss = loss + ( (self.y[j] * np.dot(self.weights, obj)) -
np.log(1 + np.exp( np.dot(self.weights, obj) )) )
            #display(self.weights)
        return loss
    def dL(self):
        dloss = 0
        for j, obj in enumerate(self.X): #for each row in X
            #self.weights = map(float,self.weights)
            #obj = map(float,obj)
            dloss = dloss + ( obj * (self.y[j] - (
(np.exp(np.dot(self.weights,obj)))/(1 + np.exp(np.dot(self.weights,obj)))
) ) )
        return dloss
    def gradientAscent(self, plot=False): #trains model, returns loss
        losses = list()
        for i in range(self.iterations):
            #display(self.weights - (self.learnRate * self.shrinkage *
self.weights))
            self.weights = self.weights - (self.learnRate * self.shrinkage *
self.weights) + ((self.learnRate/self.X.shape[1])* self.dL())
            #display(self.weights)
            losses.append(self.L())

        if (plot == True):
            itrs = range(self.iterations)
            plt.plot(itrs, losses , c='b')
            plt.xlabel("Iteration")
            plt.ylabel("Loss")

```

```

        plt.title("Loss vs. Iteration, lr = {}".format(self.learnRate))
        plt.show()
    return self.L()
def error(self, X, y):
    ret = 0
    for j, obj in enumerate(X):
        hyp = 1 / (1 + np.exp(-1 * np.dot(self.weights, obj)))
        pred = -1
        if (hyp > .5):
            pred = 1
        if (pred != y[j]):
            ret = ret + 1
    return ret / X.shape[0]

def read_csv(filename):
    with open(filename, newline='') as f_input:
        return [list(map(float, row)) for row in csv.reader(f_input)]

trainX = pd.read_csv('dexter_train.data')

trainX = np.array(trainX)
trainX = trainX.astype(str)
#trainX = trainX[:, 0:5000]
testX = pd.read_csv("dexter_valid.data")
testX = np.array(testX)
#testX = testX[:, 0:5000]
trainy = pd.read_csv("dexter_train.labels")
trainy = np.array(trainy)
testy = pd.read_csv("dexter_valid.labels")
testy = np.array(testy)

trainX_l = []
testX_l = []

for i in range(0,299):
    trainX_d = np.zeros(20000)
    #display(len(trainX_d))
    testX_d = np.zeros([20000])
    #trainX[i,0] = trainX[i,0].split(' ')

```

```

d1 = trainX[i,0].split(' ')
for j in range(len(d1)):
    q = d1[j].split(':')
    if len(q) > 1:
        q1 = int(q[0])
        q2 = int(q[1])

        trainX_d[int(q1)] = q2
trainX_l.append(trainX_d)

d2 = testX[i,0].split(' ')
for j in range(len(d2)):
    q = d2[j].split(':')
    if len(q) > 1:
        q1 = int(q[0])
        q2 = int(q[1])

        testX_d[q1] = q2

testX_l.append(testX_d)

trainX_l = np.array(trainX_l)
testX_l = np.array(testX_l)
trainX_l = np.array(trainX_l).astype(int)
testX_l = np.array(testX_l).astype(int)
trainy = np.array(trainy).astype(int)
testy = np.array(testy).astype(int)

for i in range(299):
    if trainX_l[i][:].all() != 0:
        trainX_l[i][:] = (trainX_l[i][:] - np.mean(trainX_l[i][:])) /
trainX_l[i][:]
        testX_l[i][:] = (testX_l[i][:] - np.mean(trainX_l[i][:])) /
trainX_l[i][:]

clf = LiteLR(.0005, trainX_l, trainy)
loss = clf.gradientAscent(plot=True)
print(loss)

```



```
print("Train error: {} Test error: {}".format(    clf.error(trainX_1, trainy)
, clf.error(testX_1, testy)    ))
```